



Session 37

Evaluation metrics and Model selection







Table of Content

What will We Learn Today?

1. Model evaluation
2. Evaluation Metrics
3. Holdout
4. Cross Validation
5. Hyperparameter Tuning
6. Model Selection





Model evaluation

- Metrics for Performance Evaluation
 - How to evaluate the performance of a model?
- Validation techniques
 - How to obtain reliable estimates?
- Model selection
 - How to compare the relative performance among competing models?



Performance metrics for Classification



Metrics for Performance Evaluation

- Focus on the predictive capability of a model
- Confusion Matrix:

ACTUAL CLASS	PREDICTED CLASS	
	Class=Yes (1)	Class=No (0)
	Class=Yes (1) a (TP)	Class=No (0) b (FN)
Class=No (0)	c (FP)	d (TN)

a: TP (true positive)

b: FN (false negative)

c: FP (false positive)

d: TN (true negative)

$$\text{Accuracy} = \frac{a + d}{a + b + c + d} = \frac{TP + TN}{TP + TN + FP + FN}$$



Example

Example 1

```
y_true = [0,0,0,0,0,1,1,1,1,1]
y_pred = [0,0,0,0,0,1,1,1,1,1]
```

- $\text{Accuracy} = (5+5) / (5+5+0+0) = 1$

Example 2

```
y_true = [0,0,0,0,0,1,1,1,1,1]
y_pred = [1,1,0,0,0,0,1,1,1,1]
```

- $\text{Accuracy} = (3 + 4) / (3 + 4 + 2 + 1) = 0.7$

	PREDICTED CLASS		
ACTUAL CLASS		Class=0	Class=1
	Class=0	5	0
	Class=1	0	5

	PREDICTED CLASS		
ACTUAL CLASS		Class=0	Class=1
	Class=0	3	2
	Class=1	1	4



Limitation of accuracy

- Consider a 2-class problem
 - Number of Class 0 examples = 990
 - Number of Class 1 examples = 10
- If model predicts everything to be class 0, accuracy is $990/1000 = 99\%$
 - Accuracy is misleading because model does not detect any class 1 example

ACTUAL CLASS	PREDICTED CLASS	
	Class=0	Class=1
Class=0	990	0
Class=1	10	0

TN : 990

FN : 10

TP : 0

FP : 0



Cost-sensitive measures

$$\text{sensitivity (recall)} = \frac{TP}{TP + FN}$$

$$\text{precision} = \frac{TP}{TP + FP}$$

$$\text{specificity} = \frac{TN}{TN + FP}$$

$$F_1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

	PREDICTED CLASS		
		Class=Yes	Class=No
ACTUAL CLASS	Class=Yes	(TP)	(FN)
	Class=No	(FP)	(TN)

- True positive rate (TPR) = sensitivity or recall
- True negative rate (TNR) = specificity
- Recall = How good a model is at detecting the positives
- Precision = What proportion of positive identifications was actually correct?
- F1 score = conveys the balance between the precision and the recall





Example

Example

```
y_true = [0,0,0,0,0,0,0,0,1,1]
y_pred = [1,1,0,0,0,0,0,0,0,1]
```

	PREDICTED CLASS		
		Class=0	Class=1
	Class=0	6	2
ACTUAL CLASS	Class=1	1	1

	PREDICTED CLASS		
		Class=0	Class=1
	Class=0	TN	FP
ACTUAL CLASS	Class=1	FN	TP


- Accuracy = $(TP + TN) / (TP + TN + FP + FN) = (1 + 6) / (1 + 6 + 2 + 1) = 7 / 10 = 0.7$
- Precision = $TP / (TP + FP) = 1 / (1 + 2) = 1 / 3 = 0.33$
- Recall = $TP / (TP + FN) = 1 / (1 + 1) = 1 / 2 = 0.5$
- Specificity = $TN / (TN + FP) = 6 / (6 + 2) = 6 / 8 = 0.75$
- F1 = $2 * (precision * recall) / (precision + recall) = 2 * (0.33 * 0.5) / (0.33 + 0.5) = 3.3 / 0.83 = 0.4$





Read the dataset

<https://www.kaggle.com/fedesoriano/stroke-prediction-dataset>



```
import pandas as pd
df = pd.read_csv('https://raw.githubusercontent.com/ganjar87/data_science_practice/main/healthcare-dataset-stroke-data.csv', delimiter=',')
df.head()
#df.columns
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	stroke
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	formerly smoked	1
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	NaN	never smoked	1
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92	32.5	never smoked	1
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23	34.4	smokes	1
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	174.12	24.0	never smoked	1



Evaluation metrics


```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
#scaling
scaler = StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)

model=DecisionTreeClassifier(random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

print('Accuracy : ',accuracy_score(y_test, y_pred))
print('Precision : ',precision_score(y_test, y_pred, average='binary'))
print('Recall/ sensitivity : ',recall_score(y_test, y_pred, average='binary'))
sens, spec, sup = sensitivity_specificity_support(y_test,y_pred, average='binary')
print('Specificity : ',spec)
print('F1 Score : ',f1_score(y_test, y_pred, average='binary'))

print('Confusion matrix :')
print(confusion_matrix(y_test, y_pred))
print('Plot Confusion Matrix :')
plot_confusion_matrix(model, X_test, y_test, cmap=plt.cm.Blues)
plt.show()
```

Accuracy : 0.9060665362035225
Precision : 0.19101123595505617
Recall/ sensitivity : 0.19101123595505617
Specificity : 0.9501385041551247
F1 Score : 0.1910112359550562
Confusion matrix :
[[1372 72]
 [72 17]]
Plot Confusion Matrix :





Performance metrics for Regression and Forecasting





Evaluation Metrics

- Pearson correlation coefficient (r) = measures the strength and the direction of a linear relationship between two variables. (-1 to 1)
- Coefficient determination (r^2 or r square) = gives the proportion of the variance (fluctuation) of one variable that is predictable from the other variable. (0 to 1)
- Root mean square error (RMSE) = the standard deviation of the residuals (prediction errors)

Performance Metric	Formula
Root Mean Square Error (RMSE)	$\sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$
Pearson correlation coefficient (r)	$\frac{\sum_{i=1}^n (y_i - \bar{y}_i)(\hat{y}_i - \bar{\hat{y}}_i)}{\sqrt{\sum_{i=1}^n (y_i - \bar{y}_i)^2} \sqrt{\sum_{i=1}^n (\hat{y}_i - \bar{\hat{y}}_i)^2}}$

MAE
$$\frac{1}{n} \sum_{j=1}^n |y_j - \hat{y}_j|$$





Example

- House Sales in King County, USA.
- This dataset contains house sale prices for King County, which includes Seattle. It includes homes sold between May 2014 and May 2015.
- Source : <https://www.kaggle.com/harlfoxem/housesalesprediction>

price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	yr_built
221900.0	3	1.00	1180	5650	1.0	0	0	3	7	1180	0	1955
538000.0	3	2.25	2570	7242	2.0	0	0	3	7	2170	400	1951
180000.0	2	1.00	770	10000	1.0	0	0	3	6	770	0	1933
604000.0	4	3.00	1960	5000	1.0	0	0	5	7	1050	910	1965
510000.0	3	2.00	1680	8080	1.0	0	0	3	8	1680	0	1987



Random Forest Regression

- We use library from sklearn

```
from sklearn.ensemble import RandomForestRegressor
import pandas as pd
from sklearn.model_selection import train_test_split

df_X = df.drop(['id', 'date', 'price'], axis=1)
df_y = df['price']
X = df_X.astype(float).values
y = df_y.astype(float).values

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
rf_reg = RandomForestRegressor()
rf_reg.fit(X_train, y_train)
print('coefficient of determination of training set')
print(rf_reg.score(X_train, y_train))
print('coefficient of determination of testing set')
print(rf_reg.score(X_test, y_test))
print('prediction')
y_pred = rf_reg.predict(X_test)
print(y_pred[:10])
print('real value')
print(y_test[:10])
```

```
coefficient of determination of training set
0.9822147484522787
coefficient of determination of testing set
0.8552880567206314
prediction
[ 381731.    883213.31 1087602.5  2096721.    694842.    251788.92
  839186.05  624797.99  412075.97  543389.27]
real value
[ 365000.  865000. 1038000. 1490000.  711000.  211000.  790000.  680000.
  384500.  605000.]
```




Random Forest Regression

- Calculate model performance

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
from sklearn.metrics import mean_absolute_error
from scipy import stats
import numpy as np
mse = mean_squared_error(y_test,y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test,y_pred)
mae = mean_absolute_error(y_test, y_pred)
#pakai scipy, cara 1
pearson_r, pval = stats.pearsonr(y_test, y_pred)
#pakai numpy, cara 2
r = np.corrcoef(y_test, y_pred)
print('rmse : ', rmse)
print('r2 :', r2)
print('mae : ', mae)
print('pearson r : ', pearson_r)
print('pearson r : ', r[0,1])
```

```
rmse : 143108.8916016844
r2 : 0.8581377925406867
mae : 73670.68631888017
pearson r : 0.9263856126095641
pearson r : 0.9263856126095641
```




Validation techniques





Validation techniques

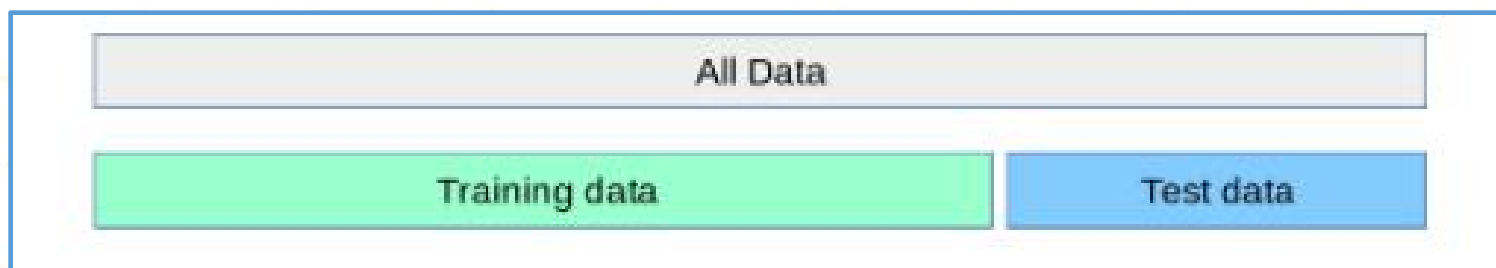
- How to obtain a reliable estimate of performance?
- Performance of a model may depend on other factors besides the learning algorithm:
 - Class distribution
 - Cost of misclassification
 - Size of training and test sets
- **Techniques**
 - Holdout
 - Cross validation



Holdout

▫ Holdout

- The data is split into two different datasets as a training and a testing dataset.
- This can be a 60/40 or 70/30 or 80/20 split.





Cross validation

- Cross validation
 - Partition data into k disjoint subsets
 - k-fold: train on k-1 partitions, test on the remaining one

Iteration 1	Test	Train	Train	Train	Train
Iteration 2	Train	Test	Train	Train	Train
Iteration 3	Train	Train	Test	Train	Train
Iteration 4	Train	Train	Train	Test	Train
Iteration 5	Train	Train	Train	Train	Test



CV version 1

▫ Using *cross_validate*

```
dt=DecisionTreeClassifier(random_state=42)
kfold = KFold(10, shuffle=True)
scores = cross_validate(dt, X, y, cv=kfold, scoring=['accuracy','precision','recall', 'f1'])

print('Accuracy : ',scores['test_accuracy'].mean())
print('Precision : ',scores['test_precision'].mean())
print('Recall/ sensitivity : ',scores['test_recall'].mean())
print('F1 : ',scores['test_f1'].mean())
```

```
Accuracy : 0.9054794520547944
Precision : 0.1156269973133051
Recall/ sensitivity : 0.1479060058947857
F1 : 0.1279194811017111
```



CV version 2

```
accs, precs, recs, specs, f1s = list(), list(), list(), list(), list()
i=1
kfold = KFold(10, shuffle=True)
for train_ix, test_ix in kfold.split(X,y):
    X_train, y_train = X[train_ix], y[train_ix]
    X_test, y_test = X[test_ix], y[test_ix]
    #print(pd.DataFrame(y_test).value_counts())
    #scaling
    scaler = StandardScaler().fit(X_train)
    X_train = scaler.transform(X_train)
    X_test = scaler.transform(X_test)

    acc, prec, rec, spec, f1 = dt(X_train, y_train, X_test, y_test)
    print('iteration ', i)
    print('accuracy ', acc, 'precision ', prec, 'recall ', rec, 'specificity ', spec, 'f1 ', f1)
    accs.append(acc)
    precs.append(prec)
    recs.append(rec)
    specs.append(spec)
    f1s.append(f1)
    print('-----')
    i = i + 1

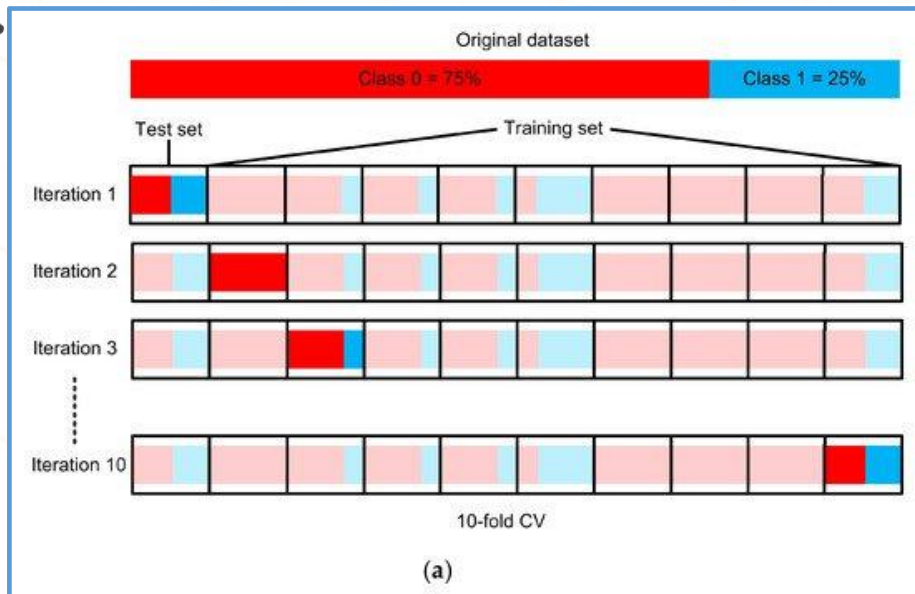
print('-----')
print('Final Accuracy: %.3f' % (np.mean(accs)))
print('Final Precision: %.3f' % (np.mean(precs)))
print('Final Specificity: %.3f' % (np.mean(specs)))
print('Final Recall: %.3f' % (np.mean(recs)))
print('Final F1: %.3f' % (np.mean(f1s)))
```



Stratified cross validation

Stratified cross validation

- Partition data into k disjoint subsets
- In stratified k -fold cross-validation, each subset is stratified so that they contain approximately the same proportion of class labels as the original dataset.
- k -fold: train on $k-1$ partitions, test on the remaining one





Stratified CV

```
accs, precs, recs, specs, f1s = list(), list(), list(), list(), list()
i=1
kfold = StratifiedKFold(10)
for train_ix, test_ix in kfold.split(X,y):
    X_train, y_train = X[train_ix], y[train_ix]
    X_test, y_test = X[test_ix], y[test_ix]
    #print(pd.DataFrame(y_test).value_counts())
    #scaling
    scaler = StandardScaler().fit(X_train)
    X_train = scaler.transform(X_train)
    X_test = scaler.transform(X_test)
    acc, prec, rec, spec, f1 = dt(X_train, y_train, X_test, y_test)
    print('iteration ', i)
    print('accuracy ', acc, 'precision ', prec, 'recall ', rec, 'specificity ', spec, 'f1 ', f1)
    accs.append(acc)
    precs.append(prec)
    recs.append(rec)
    specs.append(spec)
    f1s.append(f1)
    print('-----')
    i = i + 1

print('-----')
print('Final Accuracy: %.3f' % (np.mean(accs)))
print('Final Precision: %.3f' % (np.mean(precs)))
print('Final Specificity: %.3f' % (np.mean(specs)))
print('Final Recall: %.3f' % (np.mean(recs)))
print('Final F1: %.3f' % (np.mean(f1s)))
```




Hyperparameter tuning





Hyperparameter tuning


- Hyperparameter tuning is choosing a set of optimal hyperparameters for a learning algorithm.
- Hyperparameters = part of the input that we supply to the ML algorithm
- Parameters = found out by training the model using trainset

Hyperparameter	Parameters
NearestNeighbors(n_neighbors=5)	Linear regression
KMeans(n_clusters=2)	<pre> coefficient [-3.43081477e+04 4.03129700e+04 1.12001375e+02 9.91841247e-02 5.27154218e+03 5.43877177e+05 5.50830616e+04 2.31460673e+04 9.49081794e+04 7.22190669e+01 3.97823083e+01 -2.59441847e+03 2.19209734e+01 -5.56358731e+02 5.95216324e+05 -1.96904658e+05 1.62077488e+01 -3.30430480e-01] intercept 6641646.708113588 </pre>
DecisionTreeClassifier(max_depth=5)	
RandomForestClassifier(max_features=4)	



Grid Search

- Exhaustive search over specified hyperparameter values for an estimator.
- Steps
 - sets up a grid of hyperparameter values and for each combination,
 - trains a model and scores on the testing data.
- In this approach, every combination of hyperparameter values is tried which can be very inefficient.





Grid Search

```
dt=DecisionTreeClassifier(random_state=42)
grid_values = {'max_depth': [3, 15, 20, 40, 60], 'criterion':['gini', 'entropy' ]}
grid_dt = GridSearchCV(dt, param_grid = grid_values, scoring = 'recall', cv=10)
grid_dt.fit(X_train, y_train)
#Predict values based on new parameters
y_pred = grid_dt.predict(X_test)
#best hyperparameters
print(grid_dt.best_params_)
print(grid_dt.best_estimator_)


print('Accuracy : ', accuracy_score(y_test, y_pred))
print('Precision : ', precision_score(y_test, y_pred, average='binary'))
print('Recall/ sensitivity : ', recall_score(y_test, y_pred, average='binary'))
print('F1 : ', f1_score(y_test, y_pred, average='binary'))
sens, spec, sup = sensitivity_specificity_support(y_test, y_pred, average='binary')
print('Specificity : ', spec)
print('Confusion matrix :')
print(confusion_matrix(y_test, y_pred))
print('Plot Confusion Matrix :')
plot_confusion_matrix(grid_dt, X_test, y_test, cmap=plt.cm.Blues)
plt.show()
```





Randomized Search

- Random search is a method in which random combinations of hyperparameters are selected and used to train a model.
- The best random hyperparameter combinations are used.





Randomized Search

```
dt=DecisionTreeClassifier(random_state=42)
rand_values = {'max_depth': [3, 15, 20, 40, 60], 'criterion':['gini', 'entropy' ]}
rand_dt = RandomizedSearchCV(dt, param_distributions = rand_values, scoring = 'recall', cv=10, n_iter=3, random_state=42)
rand_dt.fit(X_train, y_train)

#Predict values based on new parameters
y_pred = rand_dt.predict(X_test)
#best hyperparameters
print(rand_dt.best_params_)
print(rand_dt.best_estimator_)
print('Accuracy : ',accuracy_score(y_test, y_pred))
print('Precision : ',precision_score(y_test, y_pred, average='binary'))
print('Recall/ sensitivity : ',recall_score(y_test, y_pred, average='binary'))
print('F1 : ',f1_score(y_test, y_pred, average='binary'))
sens, spec, sp = sensitivity_specificity_support(y_test,y_pred, average='binary')
print('Specificity : ',spec)
print('Confusion matrix :')
print(confusion_matrix(y_test, y_pred))
print('Plot Confusion Matrix :')
plot_confusion_matrix(rand_dt, X_test, y_test, cmap=plt.cm.Blues)
plt.show()
```




Model selection





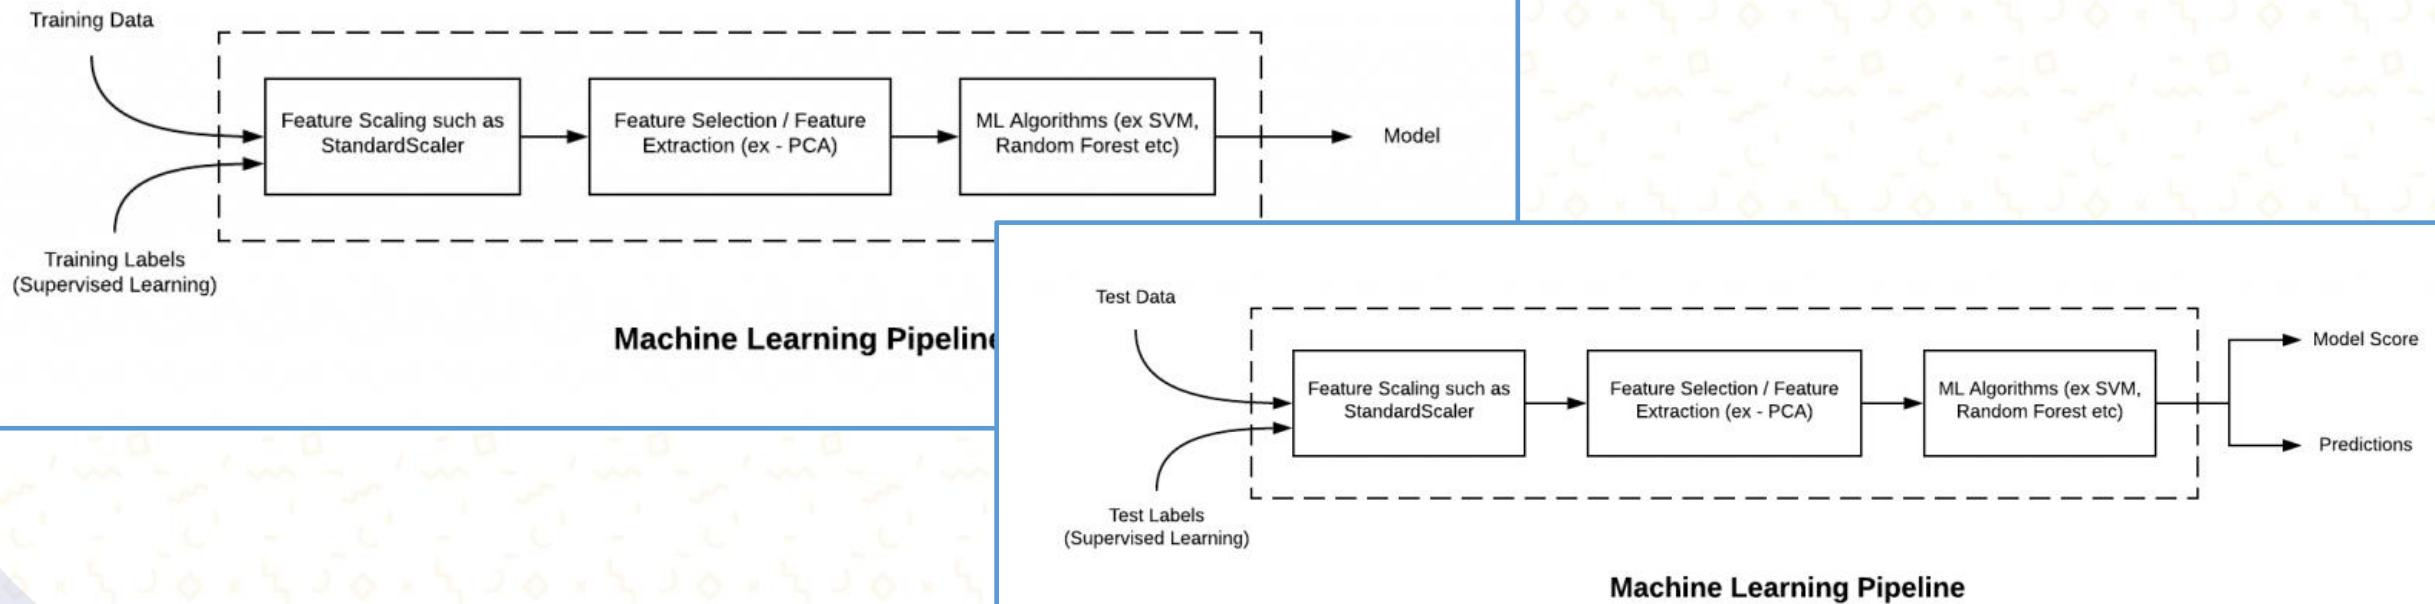
Model selection

- Model selection refers to choose the best statistical machine learning model for a particular problem.
- For this task we need to compare the relative performance between models.



ML Pipeline

- *Machine Learning (ML) pipeline* = cara meng-otomatisasi alur kerja yang diperlukan untuk menghasilkan model *machine learning*.
- Sebelumnya kita men-transformasi data untuk training dan testing set secara terpisah.
- Dengan menggunakan *Sklearn.pipeline* kita dapat meng-otomatiskan langkah-langkah ini.





Model selection

```
# Construct some pipelines
pipe_dt = Pipeline(steps=[('scaling',StandardScaler()),
                           ('classifier', DecisionTreeClassifier
                              (random_state=42, max_depth=15, criterion='gini'))])
pipe_knn = Pipeline(steps=[('scaling',StandardScaler()),
                           ('classifier', KNeighborsClassifier())])
pipe_lr = Pipeline(steps=[('scaling',StandardScaler()),
                           ('classifier', LogisticRegression(random_state=42))])
pipe_rf = Pipeline(steps=[('scaling',StandardScaler()),
                           ('classifier', RandomForestClassifier(random_state=42))])
pipe_svc = Pipeline(steps=[('scaling',StandardScaler()),
                           ('classifier', SVC(random_state=42))])
pipe_bag = Pipeline(steps=[('scaling',StandardScaler()),
                           ('classifier', BaggingClassifier(random_state=42))])
pipes = [pipe_dt, pipe_knn, pipe_lr, pipe_rf, pipe_svc, pipe_bag]
names_pipes = ['DT','KNN','LR','RF','SVM','Bagging']

for i in range(len(pipes)):
    print(names_pipes[i])
    pipes[i].fit(X_train, y_train)
    y_pred = pipes[i].predict(X_test)
    print('Accuracy : ',accuracy_score(y_test, y_pred))
    print('Precision : ',precision_score(y_test, y_pred, average='binary'))
    print('Recall/ sensitivity : ',recall_score(y_test, y_pred, average='binary'))
    print('F1 : ',f1_score(y_test, y_pred, average='binary'))
    sens, spec, sup = sensitivity_specificity_support(y_test,y_pred, average='binary')
    print('Specificity : ',spec)
    print('-----')
    print('')
```

Thank
YOU